Realeak: evidence it works (before / after)

Measured on 22 Oct 2026 on real public data: the Cabaïa catalog (2,038 products, public Shopify feed) and the latest Cabaïa listings on Vinted (FR + DE), fetched live through Keenable. Every number is traceable in the Aiven PostgreSQL database behind the live app.

1. No-data baseline vs Realeak

Question: "What share of our items listed second-hand on Vinted are brand new?"

	Answer	**Basis**
GPT-6.1 without data	About 30% (range 20 to 45%)	Its own words: low confidence, judgment-based estimate
Realeak (agent + data)	**57.8% new** (182 of 315 listings)	Measured, listing by listing

The measured value is outside the no-data range: guessing underestimates the real share by about half.

2. Keyword workflow (no AI) vs agent, on the same 315 listings

	Keyword workflow	**Realeak agent**
Listings classified into a product family	281 / 315 (89%)	315 / 315 (100%)
Exact product	145 text guesses; agreeing with the photo check on 6 of 28	47 photo-verified, the rest flagged "unknown"

Vinted titles are vague ("sac cabaia bleu"), so keywords "find" a product for half the listings but are mostly wrong. The agent names a product only when the photo supports it.

3. Time and cost

	Result
315 listings read and classified	**8.7 minutes**, fully automatic
Cost via the OpenAI API	**$0.29** (about $0.001 per listing)
Live refresh	Found **67 new listings**; after a timeout on the French site, the agent switched to the German site on its own

Limits: accuracy against human labels is not measured yet (no labels entered at submission time).